

# PostgreSQL

## Table of Contents

1. [Data Types](#)
2. [Database Operations](#)
3. [Table Operations](#)
4. [Constraints](#)
5. [Basic Queries](#)
6. [Filtering Data](#)
7. [Sorting Data](#)
8. [Joins](#)
9. [Aggregations](#)
10. [Grouping](#)
11. [Subqueries](#)
12. [Common Table Expressions \(CTEs\)](#)
13. [Window Functions](#)
14. [Indexes](#)
15. [Views](#)
16. [Functions and Stored Procedures](#)
17. [Transactions](#)
18. [Users and Privileges](#)
19. [PostgreSQL-Specific Features](#)
20. [Performance Tips](#)

## Data Types

### Numeric Types

```
-- Integers
SMALLINT          -- 2 bytes, -32768 to 32767
INTEGER           -- 4 bytes, -2147483648 to 2147483647
BIGINT            -- 8 bytes, -9223372036854775808 to
9223372036854775807
SERIAL            -- Auto-incrementing INTEGER
BIGSERIAL         -- Auto-incrementing BIGINT

-- Decimal/Floating Point
DECIMAL(p,s)      -- Exact, p=precision, s=scale
NUMERIC(p,s)      -- Exact, p=precision, s=scale
```

REAL	-- 4 bytes, inexact, 6 decimal digits precision
DOUBLE PRECISION	-- 8 bytes, inexact, 15 decimal digits precision

## Character Types

CHAR(n)	-- Fixed-length string, blank padded
VARCHAR(n)	-- Variable-length string with limit
TEXT	-- Variable unlimited length

## Boolean Type

BOOLEAN	-- true/false
---------	---------------

## Date/Time Types

DATE	-- Date only (YYYY-MM-DD)
TIME	-- Time only (HH:MM:SS)
TIMESTAMP	-- Date and time (YYYY-MM-DD HH:MM:SS)
TIMESTAMPTZ	-- Date and time with timezone
INTERVAL	-- Time interval

## Other Types

BYTEA	-- Binary data
JSON	-- JSON data
JSONB	-- Binary JSON data (more efficient)
UUID	-- Universal unique identifier
ARRAY	-- Array of values
ENUM	-- User-defined enumerated type
CIDR, INET, MACADDR	-- Network address types
POINT, LINE, POLYGON	-- Geometric types
TSQUERY, TSVECTOR	-- Text search types

## Database Operations

### Create Database

```
CREATE DATABASE database_name;  
CREATE DATABASE database_name OWNER username ENCODING 'UTF8';
```

### Drop Database

```
DROP DATABASE [IF EXISTS] database_name;
```

## List Databases

```
\l -- psql command  
SELECT datname FROM pg_database;
```

## Table Operations

### Create Table

```
CREATE TABLE table_name (  
    column1 datatype [constraints],  
    column2 datatype [constraints],  
    ...  
);  
  
-- Example  
CREATE TABLE employees (  
    employee_id SERIAL PRIMARY KEY,  
    first_name VARCHAR(50) NOT NULL,  
    last_name VARCHAR(50) NOT NULL,  
    email VARCHAR(100) UNIQUE,  
    hire_date DATE DEFAULT CURRENT_DATE,  
    salary NUMERIC(10,2)  
);
```

### Alter Table

```
-- Add column  
ALTER TABLE table_name ADD COLUMN column_name datatype [constraints];  
  
-- Drop column  
ALTER TABLE table_name DROP COLUMN column_name;  
  
-- Rename column  
ALTER TABLE table_name RENAME COLUMN old_column TO new_column;  
  
-- Change column data type  
ALTER TABLE table_name ALTER COLUMN column_name TYPE new_datatype;  
  
-- Add constraint  
ALTER TABLE table_name ADD CONSTRAINT constraint_name  
constraint_definition;
```

```
-- Drop constraint
ALTER TABLE table_name DROP CONSTRAINT constraint_name;

-- Rename table
ALTER TABLE table_name RENAME TO new_table_name;
```

## Drop Table

```
DROP TABLE [IF EXISTS] table_name [CASCADE | RESTRICT];
```

## Truncate Table (Delete all rows)

```
TRUNCATE TABLE table_name [CASCADE | RESTRICT];
```

## Copy Table Structure (without data)

```
CREATE TABLE new_table AS TABLE existing_table WITH NO DATA;
```

## Copy Table with Data

```
CREATE TABLE new_table AS SELECT * FROM existing_table;
```

## Constraints

### Primary Key

```
-- During table creation
CREATE TABLE table_name (
    id SERIAL PRIMARY KEY,
    ...
);

-- Alter existing table
ALTER TABLE table_name ADD PRIMARY KEY (column_name);
```

### Foreign Key

```
-- During table creation
CREATE TABLE orders (
    order_id SERIAL PRIMARY KEY,
    customer_id INTEGER REFERENCES customers(customer_id),
    ...
);
```

```
-- Alter existing table
ALTER TABLE orders ADD CONSTRAINT fk_customer
    FOREIGN KEY (customer_id) REFERENCES customers(customer_id)
    ON DELETE CASCADE;
```

## Unique Constraint

```
-- During table creation
CREATE TABLE users (
    username VARCHAR(50) UNIQUE,
    ...
);

-- Alter existing table
ALTER TABLE users ADD CONSTRAINT unique_email UNIQUE (email);
```

## Check Constraint

```
-- During table creation
CREATE TABLE products (
    product_id SERIAL PRIMARY KEY,
    price NUMERIC(10,2) CHECK (price > 0),
    ...
);

-- Alter existing table
ALTER TABLE products ADD CONSTRAINT positive_quantity
    CHECK (quantity >= 0);
```

## Not Null Constraint

```
-- During table creation
CREATE TABLE contacts (
    first_name VARCHAR(50) NOT NULL,
    ...
);

-- Alter existing table
ALTER TABLE contacts ALTER COLUMN last_name SET NOT NULL;
```

## Default Value

```
-- During table creation
CREATE TABLE orders (
```

```
    order_date DATE DEFAULT CURRENT_DATE,  
    ...  
);  
  
-- Alter existing table  
ALTER TABLE orders ALTER COLUMN status SET DEFAULT 'Pending';
```

## Basic Queries

### Insert Data

```
-- Insert single row  
INSERT INTO table_name (column1, column2, ...)  
VALUES (value1, value2, ...);  
  
-- Insert multiple rows  
INSERT INTO table_name (column1, column2, ...)  
VALUES  
    (value1, value2, ...),  
    (value1, value2, ...);  
  
-- Insert with returning  
INSERT INTO table_name (column1, column2, ...)  
VALUES (value1, value2, ...)  
RETURNING column1, column2;  
  
-- Insert from select  
INSERT INTO table_name (column1, column2, ...)  
SELECT column1, column2, ... FROM other_table  
WHERE condition;
```

### Update Data

```
-- Update all rows  
UPDATE table_name  
SET column1 = value1, column2 = value2  
[WHERE condition];  
  
-- Update with join  
UPDATE table_name t1  
SET column1 = t2.column1  
FROM table_name2 t2  
WHERE t1.id = t2.id;  
  
-- Update with returning  
UPDATE table_name  
SET column1 = value1
```

```
WHERE condition
RETURNING column1, column2;
```

## Delete Data

```
-- Delete rows
DELETE FROM table_name
[WHERE condition];

-- Delete with returning
DELETE FROM table_name
WHERE condition
RETURNING column1, column2;

-- Delete using joins
DELETE FROM table_name
USING another_table
WHERE table_name.id = another_table.id AND condition;
```

## Select Data

```
-- Select all columns
SELECT * FROM table_name;

-- Select specific columns
SELECT column1, column2 FROM table_name;

-- Select with column aliases
SELECT
    column1 AS alias1,
    column2 AS "Alias With Space"
FROM table_name;

-- Select with expression
SELECT
    first_name,
    last_name,
    first_name || ' ' || last_name AS full_name
FROM employees;

-- Select distinct values
SELECT DISTINCT column1 FROM table_name;

-- Select distinct combinations
SELECT DISTINCT ON (column1) column1, column2
FROM table_name
ORDER BY column1, column2;
```

```

-- Limit and offset
SELECT * FROM table_name
LIMIT 10 OFFSET 20;

-- Selecting with calculated fields
SELECT
    product_name,
    unit_price * units_in_stock AS inventory_value
FROM products;

```

## Filtering Data

### Basic WHERE Clauses

```

-- Comparison operators
SELECT * FROM table_name WHERE column_name = value;
SELECT * FROM table_name WHERE column_name > value;
SELECT * FROM table_name WHERE column_name >= value;
SELECT * FROM table_name WHERE column_name < value;
SELECT * FROM table_name WHERE column_name <= value;
SELECT * FROM table_name WHERE column_name <> value; -- Not equal
SELECT * FROM table_name WHERE column_name != value; -- Not equal

-- BETWEEN
SELECT * FROM table_name
WHERE column_name BETWEEN value1 AND value2;

-- IN operator
SELECT * FROM table_name
WHERE column_name IN (value1, value2, ...);

-- NOT IN
SELECT * FROM table_name
WHERE column_name NOT IN (value1, value2, ...);

-- LIKE (pattern matching)
SELECT * FROM table_name
WHERE column_name LIKE 'pattern%'; -- Starts with 'pattern'
SELECT * FROM table_name
WHERE column_name LIKE '%pattern'; -- Ends with 'pattern'
SELECT * FROM table_name
WHERE column_name LIKE '%pattern%'; -- Contains 'pattern'
SELECT * FROM table_name
WHERE column_name LIKE '_pattern'; -- Second to last chars are 'pattern'

-- ILIKE (case-insensitive pattern matching)
SELECT * FROM table_name

```



```
WHERE column_name ILIKE '%pattern%';
```

```
-- IS NULL / IS NOT NULL
```

```
SELECT * FROM table_name
```

```
WHERE column_name IS NULL;
```

```
SELECT * FROM table_name
```

```
WHERE column_name IS NOT NULL;
```

```
-- Logical operators
```

```
SELECT * FROM table_name
```

```
WHERE condition1 AND condition2;
```

```
SELECT * FROM table_name
```

```
WHERE condition1 OR condition2;
```

```
SELECT * FROM table_name
```

```
WHERE NOT condition;
```

## Advanced Filtering

```
-- Regex matching
```

```
SELECT * FROM table_name
```

```
WHERE column_name ~ 'regex_pattern';
```

```
-- Case-insensitive regex
```

```
SELECT * FROM table_name
```

```
WHERE column_name ~* 'regex_pattern';
```

```
-- SIMILAR TO (SQL regex)
```

```
SELECT * FROM table_name
```

```
WHERE column_name SIMILAR TO 'pattern';
```

```
-- Filtering with arrays
```

```
SELECT * FROM table_name
```

```
WHERE column_name = ANY(ARRAY[value1, value2]);
```

```
SELECT * FROM table_name
```

```
WHERE column_name @> ARRAY[value1, value2];
```

```
-- Filtering with JSON
```

```
SELECT * FROM table_name
```

```
WHERE json_column->>'property' = 'value';
```

## Sorting Data

```
-- Basic sorting
```

```
SELECT * FROM table_name
```

```
ORDER BY column1 [ASC | DESC];
```

```
-- Sorting by multiple columns
```

```
SELECT * FROM table_name
ORDER BY column1 ASC, column2 DESC;

-- Sorting with NULLS positioning
SELECT * FROM table_name
ORDER BY column1 NULLS FIRST;
SELECT * FROM table_name
ORDER BY column1 NULLS LAST;

-- Sorting with expressions
SELECT * FROM table_name
ORDER BY LOWER(column1);
SELECT * FROM table_name
ORDER BY column1 + column2;

-- Sorting by column position
SELECT column1, column2, column3 FROM table_name
ORDER BY 2 DESC; -- Sorts by column2
```

## Joins

### Inner Join

```
SELECT a.column1, b.column2
FROM table_a a
INNER JOIN table_b b ON a.key = b.key;

-- Inner join with additional conditions
SELECT a.column1, b.column2
FROM table_a a
INNER JOIN table_b b
    ON a.key = b.key AND b.column3 = 'value';
```

### Left Join (Left Outer Join)

```
SELECT a.column1, b.column2
FROM table_a a
LEFT JOIN table_b b ON a.key = b.key;
```

### Right Join (Right Outer Join)

```
SELECT a.column1, b.column2
FROM table_a a
RIGHT JOIN table_b b ON a.key = b.key;
```

## Full Join (Full Outer Join)

```
SELECT a.column1, b.column2
FROM table_a a
FULL JOIN table_b b ON a.key = b.key;
```

## Cross Join (Cartesian Product)

```
SELECT a.column1, b.column2
FROM table_a a
CROSS JOIN table_b b;
```

## Self Join

```
SELECT e.name AS employee, m.name AS manager
FROM employees e
LEFT JOIN employees m ON e.manager_id = m.employee_id;
```

## Natural Join

```
-- Joins tables using columns with the same name
SELECT a.column1, b.column2
FROM table_a a
NATURAL JOIN table_b b;
```

## Join with USING

```
-- When join columns have the same name
SELECT a.column1, b.column2
FROM table_a a
JOIN table_b b USING (common_column);
```

## Anti-Join (Finding rows in one table but not in another)

```
-- Using NOT EXISTS
SELECT a.*
FROM table_a a
WHERE NOT EXISTS (
    SELECT 1 FROM table_b b
    WHERE a.key = b.key
);

-- Using LEFT JOIN / IS NULL
```

```
SELECT a.*
FROM table_a a
LEFT JOIN table_b b ON a.key = b.key
WHERE b.key IS NULL;
```

## Aggregations

### Basic Aggregate Functions

```
SELECT
    COUNT(*) AS total_rows,
    COUNT(column1) AS non_null_values,
    COUNT(DISTINCT column1) AS unique_values,
    SUM(column1) AS total,
    AVG(column1) AS average,
    MIN(column1) AS minimum,
    MAX(column1) AS maximum
FROM table_name;
```

### Other Aggregate Functions

```
SELECT
    ARRAY_AGG(column1) AS all_values_array,
    STRING_AGG(column1, ',') AS concatenated_values,
    JSONB_AGG(column1) AS json_array,
    PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY column1) AS median,
    STDDEV(column1) AS standard_deviation,
    VARIANCE(column1) AS variance,
    MODE() WITHIN GROUP (ORDER BY column1) AS most_common_value
FROM table_name;
```

### Conditional Aggregation

```
SELECT
    SUM(CASE WHEN condition THEN column1 ELSE 0 END) AS conditional_sum,
    COUNT(CASE WHEN condition THEN 1 END) AS count_if_true,
    AVG(CASE WHEN condition THEN column1 END) AS avg_if_true
FROM table_name;
```

### FILTER Clause (PostgreSQL specific)

```
SELECT
    COUNT(*) AS total,
    COUNT(*) FILTER (WHERE condition) AS filtered_count,
```

```
SUM(column1) FILTER (WHERE condition) AS filtered_sum
FROM table_name;
```

## Grouping

### Basic GROUP BY

```
SELECT column1, COUNT(*) AS count
FROM table_name
GROUP BY column1;
```

### GROUP BY Multiple Columns

```
SELECT column1, column2, SUM(column3) AS total
FROM table_name
GROUP BY column1, column2;
```

### HAVING Clause (Filtering after grouping)

```
SELECT column1, COUNT(*) AS count
FROM table_name
GROUP BY column1
HAVING COUNT(*) > 5;
```

### GROUPING SETS

```
SELECT column1, column2, SUM(column3)
FROM table_name
GROUP BY GROUPING SETS (
    (column1, column2),
    (column1),
    (column2),
    ()
);
```

### ROLLUP (Hierarchical grouping)

```
SELECT column1, column2, SUM(column3)
FROM table_name
GROUP BY ROLLUP(column1, column2);
```

### CUBE (All possible grouping combinations)

```
SELECT column1, column2, SUM(column3)
FROM table_name
GROUP BY CUBE(column1, column2);
```

## Subqueries

### Scalar Subquery (returns a single value)

```
SELECT column1,
       (SELECT AVG(column1) FROM table_name) AS avg_value
FROM table_name
WHERE column1 > (SELECT AVG(column1) FROM table_name);
```

### Row Subquery (returns a single row)

```
SELECT *
FROM table_name
WHERE (column1, column2) = (
    SELECT column1, column2
    FROM another_table
    WHERE condition
    LIMIT 1
);
```

### Column Subquery (returns a single column)

```
SELECT *
FROM table_name
WHERE column1 IN (
    SELECT column1
    FROM another_table
    WHERE condition
);
```

### Table Subquery (returns multiple rows and columns)

```
SELECT a.column1, b.column2
FROM table_name a
JOIN (
    SELECT column1, column2
    FROM another_table
    WHERE condition
) b ON a.column1 = b.column1;
```

## EXISTS Subquery (returns boolean result)

```
SELECT *
FROM table_name a
WHERE EXISTS (
    SELECT 1
    FROM another_table b
    WHERE b.column1 = a.column1
);
```

## ANY/SOME and ALL Subqueries

```
SELECT *
FROM table_name
WHERE column1 > ANY (
    SELECT column1
    FROM another_table
    WHERE condition
);
```

```
SELECT *
FROM table_name
WHERE column1 > ALL (
    SELECT column1
    FROM another_table
    WHERE condition
);
```

## Correlated Subquery (references outer query)

```
SELECT a.column1,
       (SELECT COUNT(*)
        FROM another_table b
        WHERE b.parent_id = a.id) AS child_count
FROM table_name a;
```

## Common Table Expressions (CTEs)

### Basic CTE

```
WITH cte_name AS (
    SELECT column1, column2
    FROM table_name
    WHERE condition
)
```

```
SELECT *  
FROM cte_name  
WHERE column1 > value;
```

## Multiple CTEs

```
WITH cte1 AS (  
    SELECT column1, column2  
    FROM table1  
    WHERE condition1  
)  
cte2 AS (  
    SELECT column1, column3  
    FROM table2  
    WHERE condition2  
)  
SELECT cte1.column1, cte1.column2, cte2.column3  
FROM cte1  
JOIN cte2 ON cte1.column1 = cte2.column1;
```

## Recursive CTE

```
WITH RECURSIVE cte_name AS (  
    -- Non-recursive term (anchor)  
    SELECT column1, column2  
    FROM table_name  
    WHERE condition  
  
    UNION ALL  
  
    -- Recursive term  
    SELECT t.column1, c.column2  
    FROM table_name t  
    JOIN cte_name c ON t.parent_id = c.id  
)  
SELECT * FROM cte_name;
```

-- Example: Generate series of dates

```
WITH RECURSIVE date_series AS (  
    SELECT '2023-01-01'::date AS date  
  
    UNION ALL  
  
    SELECT date + 1  
    FROM date_series  
    WHERE date < '2023-01-31'
```



```
)  
SELECT * FROM date_series;
```

## Materialized CTE

```
-- Only in PostgreSQL 12+  
WITH cte_name AS MATERIALIZED (  
    SELECT column1, column2  
    FROM table_name  
    WHERE condition  
)  
SELECT * FROM cte_name;
```

## Window Functions

### Basic Window Function Syntax

```
SELECT  
    column1,  
    column2,  
    window_function() OVER (  
        [PARTITION BY partition_expression]  
        [ORDER BY sort_expression [ASC | DESC] [NULLS FIRST | NULLS LAST]]  
        [frame_clause]  
    ) AS alias  
FROM table_name;
```

## Ranking Functions

```
SELECT  
    column1,  
    ROW_NUMBER() OVER (PARTITION BY column2 ORDER BY column3) AS row_num,  
    RANK() OVER (PARTITION BY column2 ORDER BY column3) AS rank,  
    DENSE_RANK() OVER (PARTITION BY column2 ORDER BY column3) AS  
dense_rank,  
    PERCENT_RANK() OVER (PARTITION BY column2 ORDER BY column3) AS  
percent_rank,  
    CUME_DIST() OVER (PARTITION BY column2 ORDER BY column3) AS cume_dist,  
    NTILE(4) OVER (PARTITION BY column2 ORDER BY column3) AS quartile  
FROM table_name;
```

## Aggregate Window Functions

```
SELECT  
    column1,
```

```

column2,
SUM(column3) OVER (PARTITION BY column1) AS total_by_column1,
AVG(column3) OVER (PARTITION BY column1) AS avg_by_column1,
COUNT(*) OVER (PARTITION BY column1) AS count_by_column1,
MAX(column3) OVER (PARTITION BY column1) AS max_by_column1,
MIN(column3) OVER (PARTITION BY column1) AS min_by_column1
FROM table_name;

```

## Value Window Functions

```

SELECT
    column1,
    column2,
    FIRST_VALUE(column2) OVER (PARTITION BY column1 ORDER BY column3) AS
first_value,
    LAST_VALUE(column2) OVER (
        PARTITION BY column1
        ORDER BY column3
        RANGE BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING
    ) AS last_value,
    LEAD(column2) OVER (PARTITION BY column1 ORDER BY column3) AS
next_value,
    LEAD(column2, 2, 'N/A') OVER (PARTITION BY column1 ORDER BY column3)
AS second_next_value,
    LAG(column2) OVER (PARTITION BY column1 ORDER BY column3) AS
previous_value,
    LAG(column2, 2, 'N/A') OVER (PARTITION BY column1 ORDER BY column3) AS
second_previous_value
FROM table_name;

```

## Window Frames

```

SELECT
    column1,
    column2,
    SUM(column3) OVER (
        PARTITION BY column1
        ORDER BY column2
        ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING
    ) AS moving_sum,
    AVG(column3) OVER (
        PARTITION BY column1
        ORDER BY column2
        ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
    ) AS running_avg
FROM table_name;

```

# Indexes

## Create Index

```
-- Basic index
CREATE INDEX index_name ON table_name (column_name);

-- Multi-column index
CREATE INDEX index_name ON table_name (column1, column2);

-- Unique index
CREATE UNIQUE INDEX index_name ON table_name (column_name);

-- Partial index
CREATE INDEX index_name ON table_name (column_name)
WHERE condition;

-- Expression index
CREATE INDEX index_name ON table_name (LOWER(column_name));

-- GIN index (for arrays, jsonb, etc.)
CREATE INDEX index_name ON table_name USING GIN (column_name);

-- GIST index (for geometric data types, etc.)
CREATE INDEX index_name ON table_name USING GIST (column_name);

-- BRIN index (for large tables with sorted values)
CREATE INDEX index_name ON table_name USING BRIN (column_name);

-- Hash index (equality operations only)
CREATE INDEX index_name ON table_name USING HASH (column_name);
```

## Alter Index

```
-- Rename index
ALTER INDEX index_name RENAME TO new_index_name;

-- Set tablespace
ALTER INDEX index_name SET TABLESPACE tablespace_name;
```

## Drop Index

```
DROP INDEX [IF EXISTS] index_name [CASCADE | RESTRICT];
```

## Reindex

```
-- Rebuild single index
REINDEX INDEX index_name;

-- Rebuild all indexes on a table
REINDEX TABLE table_name;

-- Rebuild all indexes in a database
REINDEX DATABASE database_name;
```

## Views

### Create View

```
-- Basic view
CREATE VIEW view_name AS
SELECT column1, column2
FROM table_name
WHERE condition;

-- Updatable view
CREATE VIEW view_name AS
SELECT column1, column2
FROM table_name
WHERE condition
WITH [CASCADED | LOCAL] CHECK OPTION;

-- Materialized view (saved result set)
CREATE MATERIALIZED VIEW view_name AS
SELECT column1, column2
FROM table_name
WHERE condition;
```

### Alter View

```
-- Rename view
ALTER VIEW view_name RENAME TO new_view_name;

-- Change owner
ALTER VIEW view_name OWNER TO new_owner;
```

### Drop View

```
-- Regular view
DROP VIEW [IF EXISTS] view_name [CASCADE | RESTRICT];
```

```
-- Materialized view
DROP MATERIALIZED VIEW [IF EXISTS] view_name [CASCADE | RESTRICT];
```

## Refresh Materialized View

```
-- Complete refresh
REFRESH MATERIALIZED VIEW view_name;

-- Concurrent refresh (doesn't block reads)
REFRESH MATERIALIZED VIEW CONCURRENTLY view_name;
```

## Functions and Stored Procedures

### Create Function

```
CREATE [OR REPLACE] FUNCTION function_name(param1 type, param2 type)
RETURNS return_type
LANGUAGE language_name
[IMMUTABLE | STABLE | VOLATILE]
[PARALLEL SAFE | PARALLEL RESTRICTED | PARALLEL UNSAFE]
AS $$
    -- Function body
$$;

-- Example PL/pgSQL function
CREATE OR REPLACE FUNCTION get_employee_salary(emp_id integer)
RETURNS numeric
LANGUAGE plpgsql
AS $$
DECLARE
    emp_salary numeric;
BEGIN
    SELECT salary INTO emp_salary
    FROM employees
    WHERE employee_id = emp_id;

    RETURN emp_salary;
END;
$$;
```

### Create Procedure (PostgreSQL 11+)

```
CREATE [OR REPLACE] PROCEDURE procedure_name(param1 type, param2 type)
LANGUAGE language_name
AS $$
    -- Procedure body
```

```

$$;

-- Example PL/pgSQL procedure
CREATE OR REPLACE PROCEDURE update_employee_salary(
    emp_id integer,
    new_salary numeric
)
LANGUAGE plpgsql
AS $$
BEGIN
    UPDATE employees
    SET salary = new_salary
    WHERE employee_id = emp_id;

    COMMIT;
END;
$$;

```

## Alter Function/Procedure

```

-- Rename function
ALTER FUNCTION function_name(param_types) RENAME TO new_name;

-- Change owner
ALTER FUNCTION function_name(param_types) OWNER TO new_owner;

-- Rename procedure (PostgreSQL 11+)
ALTER PROCEDURE procedure_name(param_types) RENAME TO new_name;

```

## Drop Function/Procedure

```

-- Drop function
DROP FUNCTION [IF EXISTS] function_name(param_types) [CASCADE | RESTRICT];

-- Drop procedure (PostgreSQL 11+)
DROP PROCEDURE [IF EXISTS] procedure_name(param_types) [CASCADE | RESTRICT];

```

## Call Procedure (PostgreSQL 11+)

```

CALL procedure_name(arg1, arg2);

```

## Transactions

### Basic Transaction Control

```
-- Begin a transaction
BEGIN;
-- or
BEGIN TRANSACTION;
-- or
BEGIN WORK;

-- Execute statements
INSERT INTO table_name VALUES (1, 'value');
UPDATE another_table SET column = 'new value' WHERE id = 1;

-- Commit changes
COMMIT;
-- or
COMMIT TRANSACTION;
-- or
COMMIT WORK;

-- Rollback (undo) changes
ROLLBACK;
-- or
ROLLBACK TRANSACTION;
-- or
ROLLBACK WORK;
```

## Savepoints

```
-- Begin transaction
BEGIN;

-- Execute statements
INSERT INTO table_name VALUES (1, 'value');

-- Create a savepoint
SAVEPOINT my_savepoint;

-- Execute more statements
UPDATE table_name SET column = 'new value' WHERE id = 1;

-- Rollback to savepoint (undo updates but keep inserts)
ROLLBACK TO SAVEPOINT my_savepoint;

-- Release savepoint (remove it)
RELEASE SAVEPOINT my_savepoint;

-- Commit remaining changes
COMMIT;
```

# Transaction Isolation Levels

```
-- Set transaction isolation level
BEGIN;
SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED;
-- or
SET TRANSACTION ISOLATION LEVEL READ COMMITTED;
-- or
SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;
-- or
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;

-- Execute statements
SELECT * FROM table_name;

-- Commit
COMMIT;
```

## Transaction Modes

```
-- Read-only transaction
BEGIN;
SET TRANSACTION READ ONLY;

-- Read-write transaction (default)
BEGIN;
SET TRANSACTION READ WRITE;

-- Deferrable transaction
BEGIN;
SET TRANSACTION DEFERRABLE;
```

# Users and Privileges

## Create User/Role

```
-- Create user
CREATE USER username WITH PASSWORD 'password';

-- Create role (can be assigned to multiple users)
CREATE ROLE role_name;

-- Create role with login capability
CREATE ROLE role_name WITH LOGIN PASSWORD 'password';
```



```
-- Create superuser
CREATE ROLE role_name WITH SUPERUSER LOGIN PASSWORD 'password';
```

## Alter User/Role

```
-- Change password
ALTER USER username WITH PASSWORD 'new_password';

-- Add attributes
ALTER ROLE role_name WITH CREATEDB CREATEROLE;

-- Remove attributes
ALTER ROLE role_name WITH NOCREATEDB NOCREATEROLE;
```

## Drop User/Role

```
DROP USER [IF EXISTS] username;
DROP ROLE [IF EXISTS] role_name;
```

## Grant Privileges

```
-- Grant table privileges
GRANT SELECT, INSERT, UPDATE, DELETE ON table_name TO role_name;

-- Grant all privileges on table
GRANT ALL PRIVILEGES ON table_name TO role_name;

-- Grant column-level privileges
GRANT SELECT (column1, column2), UPDATE (column1) ON table_name TO
role_name;

-- Grant database privileges
GRANT CREATE ON DATABASE database_name TO role_name;

-- Grant schema privileges
GRANT USAGE ON SCHEMA schema_name TO role_name;

-- Grant with grant option (recipient can grant to others)
GRANT SELECT ON table_name TO role_name WITH GRANT OPTION;
```

## Revoke Privileges

```
-- Revoke table privileges
REVOKE SELECT, INSERT ON table_name FROM role_name;
```

```
-- Revoke all privileges
REVOKE ALL PRIVILEGES ON table_name FROM role_name;

-- Revoke grant option only
REVOKE GRANT OPTION FOR SELECT ON table_name FROM role_name;
```

## Role Membership

```
-- Grant role to another role
GRANT role1 TO role2;

-- Revoke role from another role
REVOKE role1 FROM role2;
```

## PostgreSQL-Specific Features

### Arrays

```
-- Create array column
CREATE TABLE table_name (
    id serial PRIMARY KEY,
    tags TEXT[]
);

-- Insert arrays
INSERT INTO table_name (tags) VALUES (ARRAY['tag1', 'tag2', 'tag3']);
INSERT INTO table_name (tags) VALUES ('{"tag1", "tag2", "tag3"}');

-- Query arrays
SELECT * FROM table_name WHERE 'tag1' = ANY(tags);
SELECT * FROM table_name WHERE tags @> ARRAY['tag1', 'tag2'];
SELECT * FROM table_name WHERE tags && ARRAY['tag1', 'tag4']; -- Overlap

-- Array functions
SELECT array_length(tags, 1) FROM table_name; -- Length of dimension 1
SELECT array_to_string(tags, ', ') FROM table_name; -- Join elements
SELECT unnest(tags) FROM table_name; -- Expand array to rows
```

### JSON/JSONB

```
-- Create JSON/JSONB column
CREATE TABLE table_name (
    id serial PRIMARY KEY,
    data JSON,
    data_b JSONB -- Binary JSON, more efficient for operations
);
```

```

-- Insert JSON data
INSERT INTO table_name (data, data_b)
VALUES ('{"name": "John", "age": 30}', '{"name": "John", "age": 30}');

-- Query JSON
SELECT data->>'name' AS name FROM table_name; -- Get as text
SELECT data->'age' AS age FROM table_name;      -- Get as JSON
SELECT * FROM table_name WHERE data->>'name' = 'John';

-- JSONB operations (more efficient than JSON)
SELECT * FROM table_name WHERE data_b @> '{"name": "John"}';
SELECT * FROM table_name WHERE data_b ? 'name'; -- Has key
SELECT * FROM table_name WHERE data_b ?| ARRAY['name', 'email']; -- Has
any key
SELECT * FROM table_name WHERE data_b ?& ARRAY['name', 'age']; -- Has all
keys

-- JSONB modification
UPDATE table_name SET data_b = data_b || '{"email":
"john@example.com"}'::jsonb;
UPDATE table_name SET data_b = jsonb_set(data_b, '{age}', '31');
UPDATE table_name SET data_b = data_b - 'email'; -- Remove key

```

## Full-Text Search

```

-- Create tsvector column for full-text search
CREATE TABLE articles (
    id serial PRIMARY KEY,
    title TEXT,
    body TEXT,
    search_vector TSVECTOR
);

CREATE INDEX articles_search_idx ON articles USING GIN (search_vector);

-- Update search vector
UPDATE articles
SET search_vector =
    setweight(to_tsvector('english', COALESCE(title, '')), 'A') ||
    setweight(to_tsvector('english', COALESCE(body, '')), 'B');

-- Simple search
SELECT id, title, body
FROM articles
WHERE search_vector @@ to_tsquery('english', 'search & terms');

-- Ranking search results

```

```

SELECT id, title, body,
       ts_rank(search_vector, query) AS rank
FROM articles, to_tsquery('english', 'search & terms') query
WHERE search_vector @@ query
ORDER BY rank DESC;

-- Highlight search results
SELECT id, title,
       ts_headline('english', body, to_tsquery('english', 'search & terms'))
FROM articles
WHERE search_vector @@ to_tsquery('english', 'search & terms');

```

## Inheritance

```

-- Create parent table
CREATE TABLE cities (
    id serial PRIMARY KEY,
    name text,
    population integer
);

-- Create child table that inherits
CREATE TABLE capitals (
    country text,
    elevation integer
) INHERITS (cities);

-- Query from parent (includes child rows)
SELECT * FROM cities;

-- Query only from parent
SELECT * FROM ONLY cities;

```

## Table Partitioning

```

-- Declarative partitioning (PostgreSQL 10+)

-- Create partitioned table
CREATE TABLE measurements (
    id serial,
    logdate date,
    peaktemp integer,
    unitsales integer
) PARTITION BY RANGE (logdate);

-- Create partitions
CREATE TABLE measurements_y2022 PARTITION OF measurements

```

```
FOR VALUES FROM ('2022-01-01') TO ('2023-01-01');

CREATE TABLE measurements_y2023 PARTITION OF measurements
FOR VALUES FROM ('2023-01-01') TO ('2024-01-01');

-- Insert data (automatically routes to correct partition)
INSERT INTO measurements (logdate, peaktemp, unitsales)
VALUES ('2022-06-01', 98, 1234);
```

## Extensions

```
-- List available extensions
SELECT name, default_version, installed_version
FROM pg_available_extensions;

-- Install extension
CREATE EXTENSION extension_name;

-- Common extensions
CREATE EXTENSION postgis;           -- Spatial database
CREATE EXTENSION pg_stat_statements; -- Query statistics
CREATE EXTENSION pgcrypto;          -- Cryptographic functions
CREATE EXTENSION uuid-oss;           -- UUID generation
CREATE EXTENSION hstore;             -- Key-value store
CREATE EXTENSION pg_trgm;            -- Trigram matching for fuzzy search
```

## Performance Tips

### Explain Analyze

```
-- View query plan
EXPLAIN SELECT * FROM table_name WHERE condition;

-- Execute query and show actual times and row counts
EXPLAIN ANALYZE SELECT * FROM table_name WHERE condition;

-- Format output as JSON
EXPLAIN (FORMAT JSON) SELECT * FROM table_name WHERE condition;

-- Verbose output with more details
EXPLAIN (ANALYZE, VERBOSE) SELECT * FROM table_name WHERE condition;
```

## Table Statistics

```
-- Update statistics
ANALYZE [table_name];
```

```
-- Set statistics target for a column
ALTER TABLE table_name ALTER COLUMN column_name SET STATISTICS 1000;
```

## Optimize Queries

```
-- Use parameterized queries in your application code
-- Bad:  SELECT * FROM users WHERE name = 'John';
-- Good: SELECT * FROM users WHERE name = $1; (with parameter 'John')

-- Avoid SELECT * - specify only needed columns
-- Bad:  SELECT * FROM large_table;
-- Good: SELECT id, name FROM large_table;

-- Use EXISTS for existence check
-- Bad:  SELECT 1 FROM table WHERE EXISTS (SELECT * FROM another_table);
-- Good: SELECT 1 FROM table WHERE EXISTS (SELECT 1 FROM another_table);

-- Use COUNT(*) for counting all rows
-- Bad:  SELECT COUNT(id) FROM table;
-- Good: SELECT COUNT(*) FROM table;

-- Use appropriate indexes for JOIN, WHERE, ORDER BY
```

## Vacuuming

```
-- Reclaim space and update statistics
VACUUM [table_name];

-- More thorough vacuum (locks table)
VACUUM FULL [table_name];

-- Run vacuum and analyze in one command
VACUUM ANALYZE [table_name];
```

## Monitoring

```
-- Check activity
SELECT * FROM pg_stat_activity;

-- Identify slow queries
SELECT query, calls, total_time, rows,
       (total_time/calls) as avg_time
FROM pg_stat_statements
ORDER BY total_time DESC
LIMIT 20;
```

```
-- Check index usage
SELECT * FROM pg_stat_user_indexes;

-- Check table I/O statistics
SELECT * FROM pg_statio_user_tables;
```

## Connection Pooling

```
-- Not SQL, but configuration tips:
-- 1. Use a connection pooler like PgBouncer or Odyssey
-- 2. Configure max_connections appropriately
-- 3. Reuse connections in application code
```